



(Linear) dynamical systems

A tool for neural and behavioural data analysis

Kris Jensen, February 2025



Previously: linear regression



Assumptions:

- We have access to labelled data
- Data points are i.i.d. in time


Previously: PCA



Assumptions:




- ~~We have access to labelled data~~
- Data points are i.i.d. in time

What if this is not true?




Example 1: behavioural tracking from noisy measurements

What if this is not true?




Example 2: neurons as a noisy readout of a low-dimensional system

What if this is not true?



Example 3: a simple 'brain-computer interface'

PCA as a *probabilistic generative model*




Optimize $\mathbf{C} = \underset{\mathbf{C}}{\operatorname{argmax}} p(\mathbf{x}; \mathbf{C}) = \int p(\mathbf{x}|\mathbf{z}; \mathbf{C})p(\mathbf{z})d\mathbf{z}$

Compute $p(\mathbf{z}|\mathbf{x}; \mathbf{C})$

Upside: easy to extend & uncertainty estimates

The Kalman filter



$$p(\mathbf{z}_t) = \mathcal{N}(\mathbf{A}\mathbf{z}_{t-1}, \mathbf{Q})$$

$$p(\mathbf{x}_t|\mathbf{z}_t) = \mathcal{N}(\mathbf{C}\mathbf{z}_t, \mathbf{R})$$

The Kalman filter – inference (1d)

$z_t \sim \mathcal{N}(az_{t-1}, q^2); \quad x_t \sim \mathcal{N}(cz_t, r^2).$ Our goal: estimate $p(z_t|x_{1:t}) = \mathcal{N}(\mu_t, \sigma_t^2)$

- $p(z_t|x_{1:t}) \propto p(x_t|z_t, x_{1:t-1})p(z_t|x_{1:t-1}) = p(x_t|z_t)p(z_t|x_{1:t-1})$ (prediction & update)
- $p(z_t|x_{1:t-1}) = \int p(z_t, z_{t-1}|x_{1:t-1})dz_{t-1} = \int p(z_t|z_{t-1})p(z_{t-1}|x_{1:t-1})dz_{t-1}$
- $p(z_t|z_{t-1}) = \mathcal{N}(az_{t-1}, q^2); \quad p(z_{t-1}|x_{1:t-1}) = \mathcal{N}(\mu_{t-1}, \sigma_{t-1}^2)$
- $p(z_t|x_{1:t-1}) = \mathcal{N}(\mu_{t|t-1}, \sigma_{t|t-1}^2)$
 - $\mu_{t|t-1} = \langle a(\mu_{t-1} + \sigma_{t-1}\epsilon_1) + q\epsilon_2 \rangle_{\epsilon_1, \epsilon_2} = a\mu_{t-1}, \quad \text{where } \epsilon \sim \mathcal{N}(0,1)$
 - $\sigma_{t|t-1}^2 = \langle (a(\mu_{t-1} + \sigma_{t-1}\epsilon_1 - \mu_{t-1}) + q\epsilon_2)^2 \rangle_{\epsilon_1, \epsilon_2} = a^2\sigma_{t-1}^2 + q^2$
- $p(x_t|z_t) = \mathcal{N}(cz_t, r^2)$
 - $\exp[-r^{-2}(x_t - cz_t)^2] = \exp[-(r/c)^{-2}(z_t - x_t/c)^2]$
- $p(z_t|x_{1:t}) = \mathcal{N}(\mu_t, \sigma_t^2)$
 - $\sigma_t^2 = \left((c/r)^2 + \sigma_{t|t-1}^{-2} \right)^{-1}$
 - $\mu_t = \sigma_t^2 \left(\left(\frac{c}{r} \right)^2 \frac{x_t}{c} + \sigma_{t|t-1}^{-2} \mu_{t|t-1} \right)$

The Kalman filter – inference (general)

The equations for higher-dimensional systems can be derived similarly.

$$\mathbf{z}_t \sim \mathcal{N}(\mathbf{A}\mathbf{z}_{t-1}, \mathbf{Q}); \quad \mathbf{x}_t \sim \mathcal{N}(\mathbf{C}\mathbf{z}_t, \mathbf{R})$$


- $p(\mathbf{z}_t | \mathbf{x}_{1:t-1}) = \mathcal{N}(\boldsymbol{\mu}_{t|t-1}, \boldsymbol{\Sigma}_{t|t-1})$
 - $\boldsymbol{\mu}_{t|t-1} = \mathbf{A}\boldsymbol{\mu}_{t-1}$
 - $\boldsymbol{\Sigma}_{t|t-1} = \mathbf{A}\boldsymbol{\Sigma}_{t-1}\mathbf{A}^T + \mathbf{Q}$
- $p(\mathbf{x}_t | \mathbf{z}_t) = \mathcal{N}(\mathbf{C}\mathbf{z}_t, \mathbf{R})$
- $p(\mathbf{z}_t | \mathbf{x}_{1:t}) = \mathcal{N}(\boldsymbol{\mu}_t, \boldsymbol{\Sigma}_t)$
 - $\boldsymbol{\Sigma}_t = (\mathbf{C}^T \mathbf{R}^{-1} \mathbf{C} + \boldsymbol{\Sigma}_{t|t-1}^{-1})^{-1}$
 - $\boldsymbol{\mu}_t = \boldsymbol{\Sigma}_t (\mathbf{C}^T \mathbf{R}^{-1} \mathbf{x}_t + \boldsymbol{\Sigma}_{t|t-1}^{-1} \boldsymbol{\mu}_{t|t-1})$

Kalman *smoothing*

$$\begin{aligned} p(z_t|x_{1:T}) &= \int p(z_t, z_{t+1}|x_{1:T})dz_{t+1} \\ &= \int p(z_t|z_{t+1}, x_{1:T})p(z_{t+1}|x_{1:T})dz_{t+1} \\ &= \int p(z_t|z_{t+1}, x_{1:t})p(z_{t+1}|x_{1:T})dz_{t+1} \\ &\propto p(z_t|x_{1:t}) \int p(z_{t+1}|z_t)p(z_{t+1}|x_{1:T})dz_{t+1} \end{aligned}$$

- This recursion expresses $p(z_t|x_{1:T}) = \mathcal{N}(v_t, \eta_t^2)$ as a function of (v_{t+1}, η_{t+1}^2) , our ‘filtering distribution’ $z_t|x_{1:t} \sim \mathcal{N}(\mu_t, \sigma_t^2)$ and dynamics $z_{t+1}|z_t \sim \mathcal{N}(az_t, q^2)$.
- Upside: we get a better estimate of our latents by using both past and future data.
- Downside: we have to wait until the experiment finishes before we can perform the computation.
- This is useful for post-hoc data processing, but not for online decoding/tracking.
- An intermediate approach computes $p(z_t|x_{1:t+n})$ for some choice of n .
- For more details, see *Roweis & Ghahramani (1999)*.

Where do the parameters come from???




We might know the ‘true’ dynamics

(https://github.com/joacorapela/lds_python/blob/master/docs/tracking/tracking.pdf)

We can *maximize the likelihood* of the parameters on some training data:

$$L(\theta) = \sum_t \log p(\mathbf{z}_t | \mathbf{x}_{1:t}, \theta)$$

What if we don't have any 'training data'?



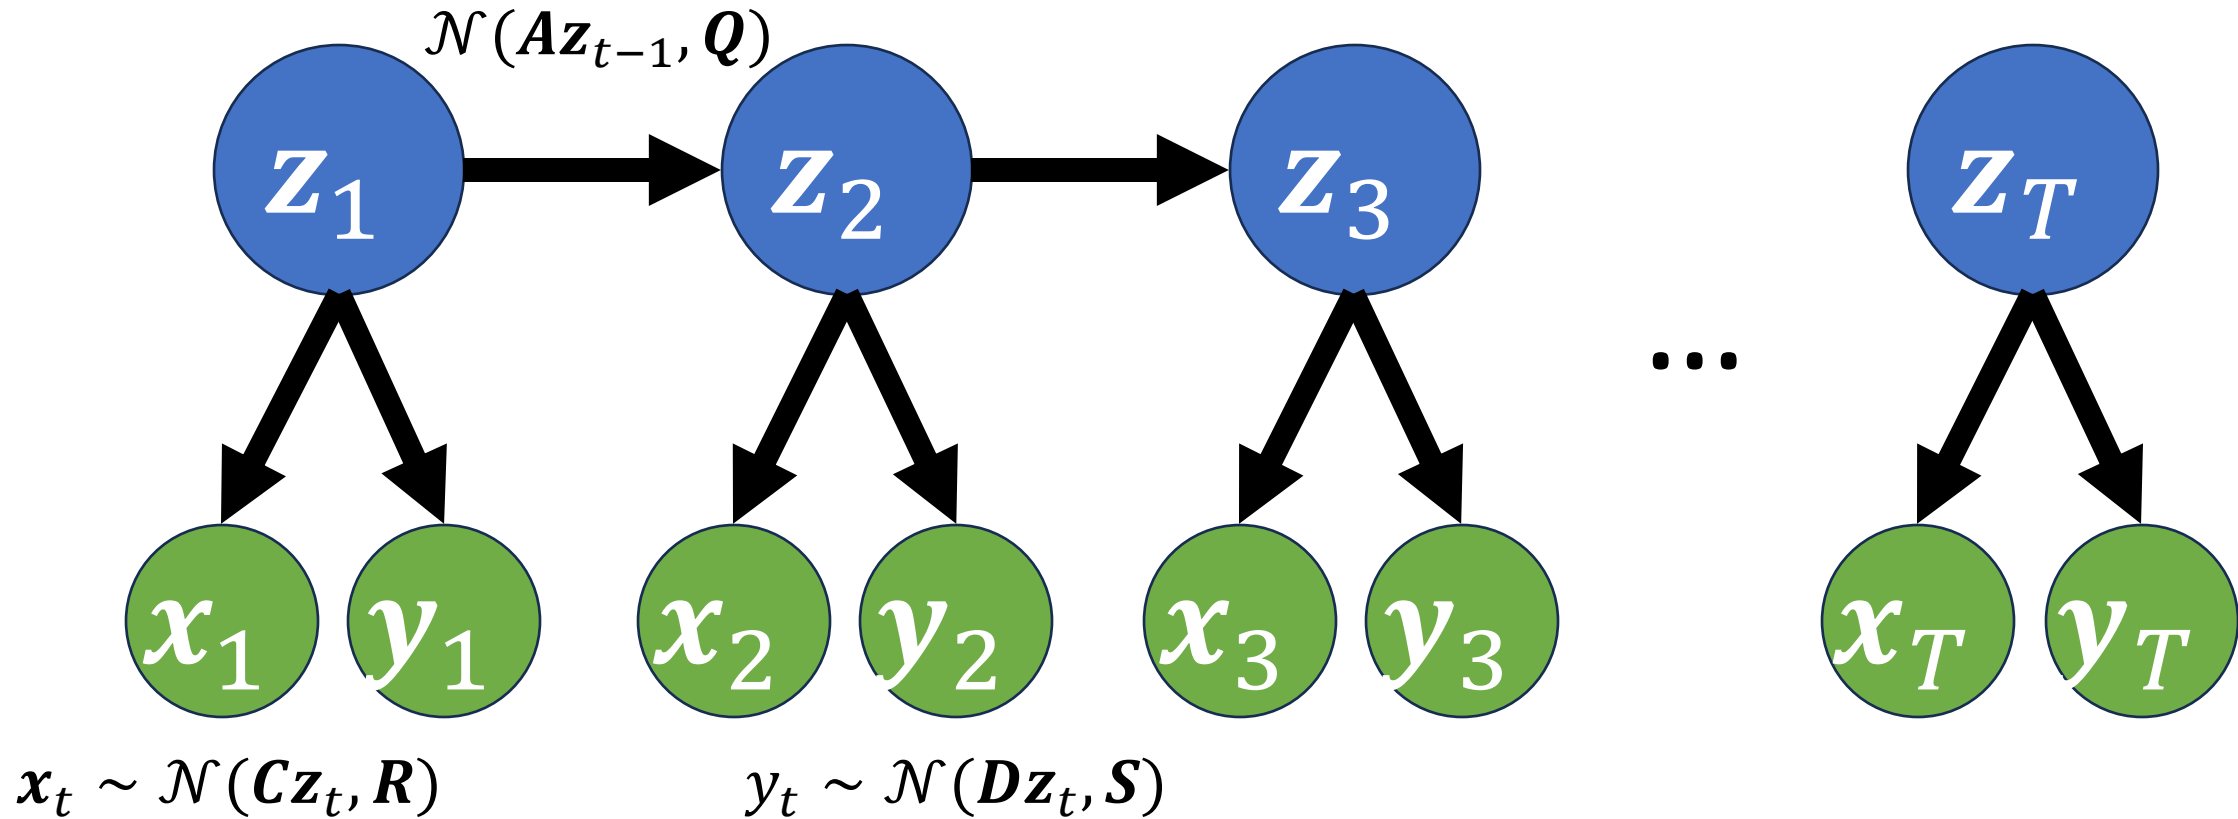
We can maximize the likelihood of the parameters on the *observations*:

$$L(\theta) = \log p(\mathbf{X}|\theta) = \log \int_{\mathbf{Z}} p(\mathbf{X}|\mathbf{Z}, \theta) p(\mathbf{Z}|\theta) d\mathbf{Z}$$

We use *expectation maximization*, alternating between inferring \mathbf{Z} and optimizing θ :

$$\theta_{i+1} \leftarrow \underset{\theta}{\operatorname{argmax}} \mathbb{E}_{\mathbf{Z} \sim p(\mathbf{Z}|\mathbf{X}; \theta_i)} [\log p(\mathbf{Z}, \mathbf{X}; \theta)]$$

What if we care about the *output* of an LDS?

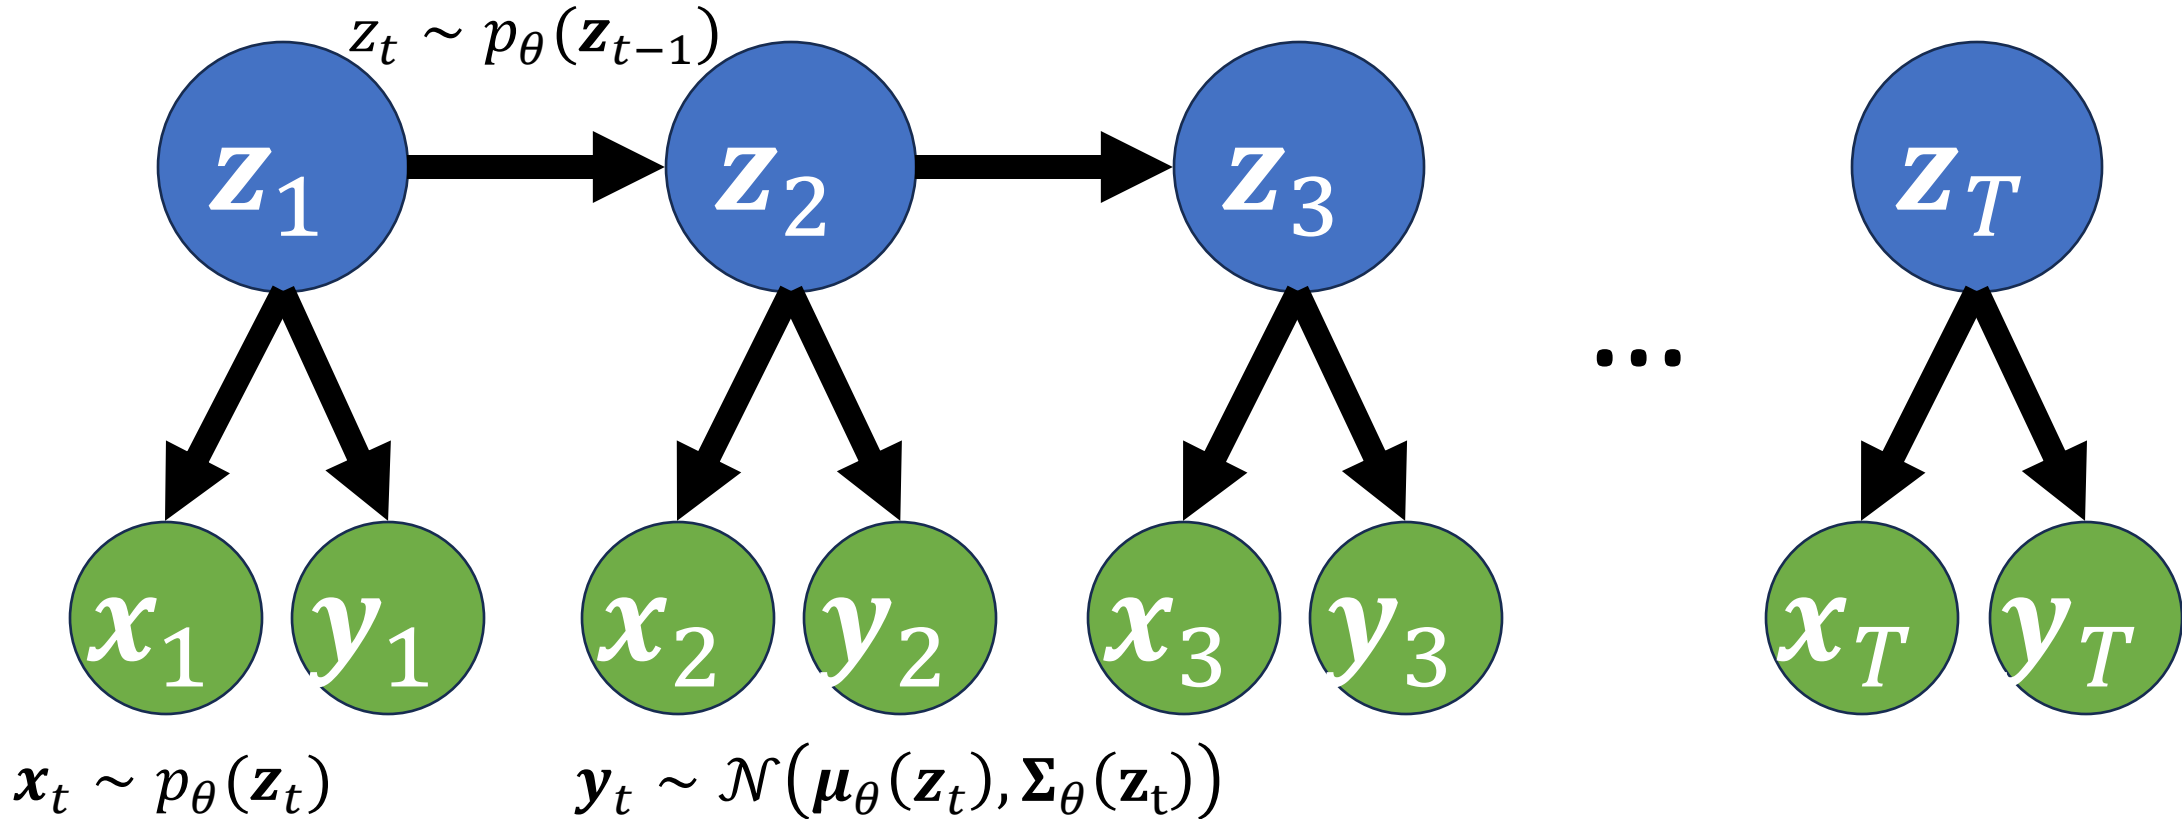


We can compute $p(\mathbf{y}_t | \mathbf{x}_{1:t}) = \int_{\mathbf{z}} p(\mathbf{y}_t | \mathbf{z}_t) p(\mathbf{z}_t | \mathbf{x}_{1:t}) d\mathbf{z}$

$p(\mathbf{z}_t | \mathbf{x}_{1:t})$ is given by our Kalman filter

$\{\mathbf{A}, \mathbf{C}, \mathbf{D}, \mathbf{Q}, \mathbf{R}, \mathbf{S}\} = \underset{\theta}{\operatorname{argmax}} \log p(\mathbf{Y}, \mathbf{X})$

Why linear and why Gaussian?

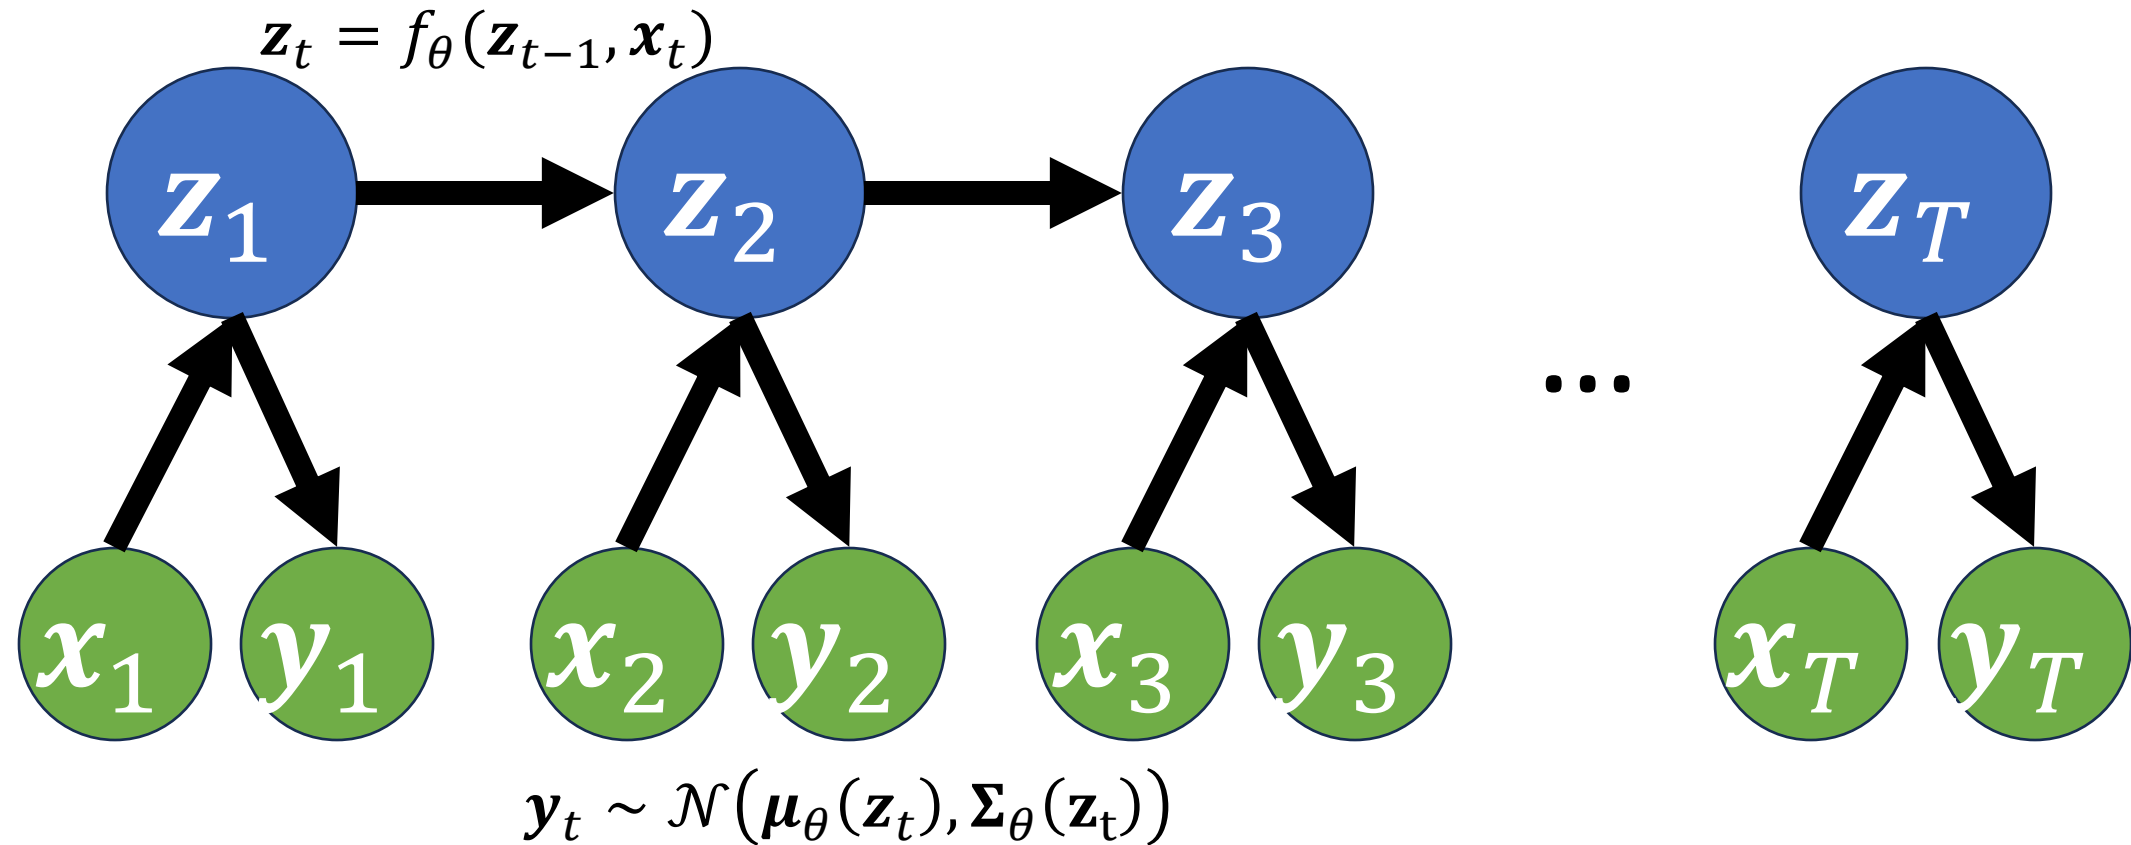


~~We can compute $p(\mathbf{y}_t | \mathbf{x}_{1:t}) = \int_{\mathbf{z}} p(\mathbf{y}_t | \mathbf{z}_t) p(\mathbf{z}_t | \mathbf{x}_{1:t}) d\mathbf{z}$~~

~~$p(\mathbf{z}_t | \mathbf{x}_{1:t})$ is given by our Kalman filter~~

$$\theta = \underset{\theta}{\operatorname{argmax}} \log p(\mathbf{Y}, \mathbf{X})$$

Why linear and why Gaussian?



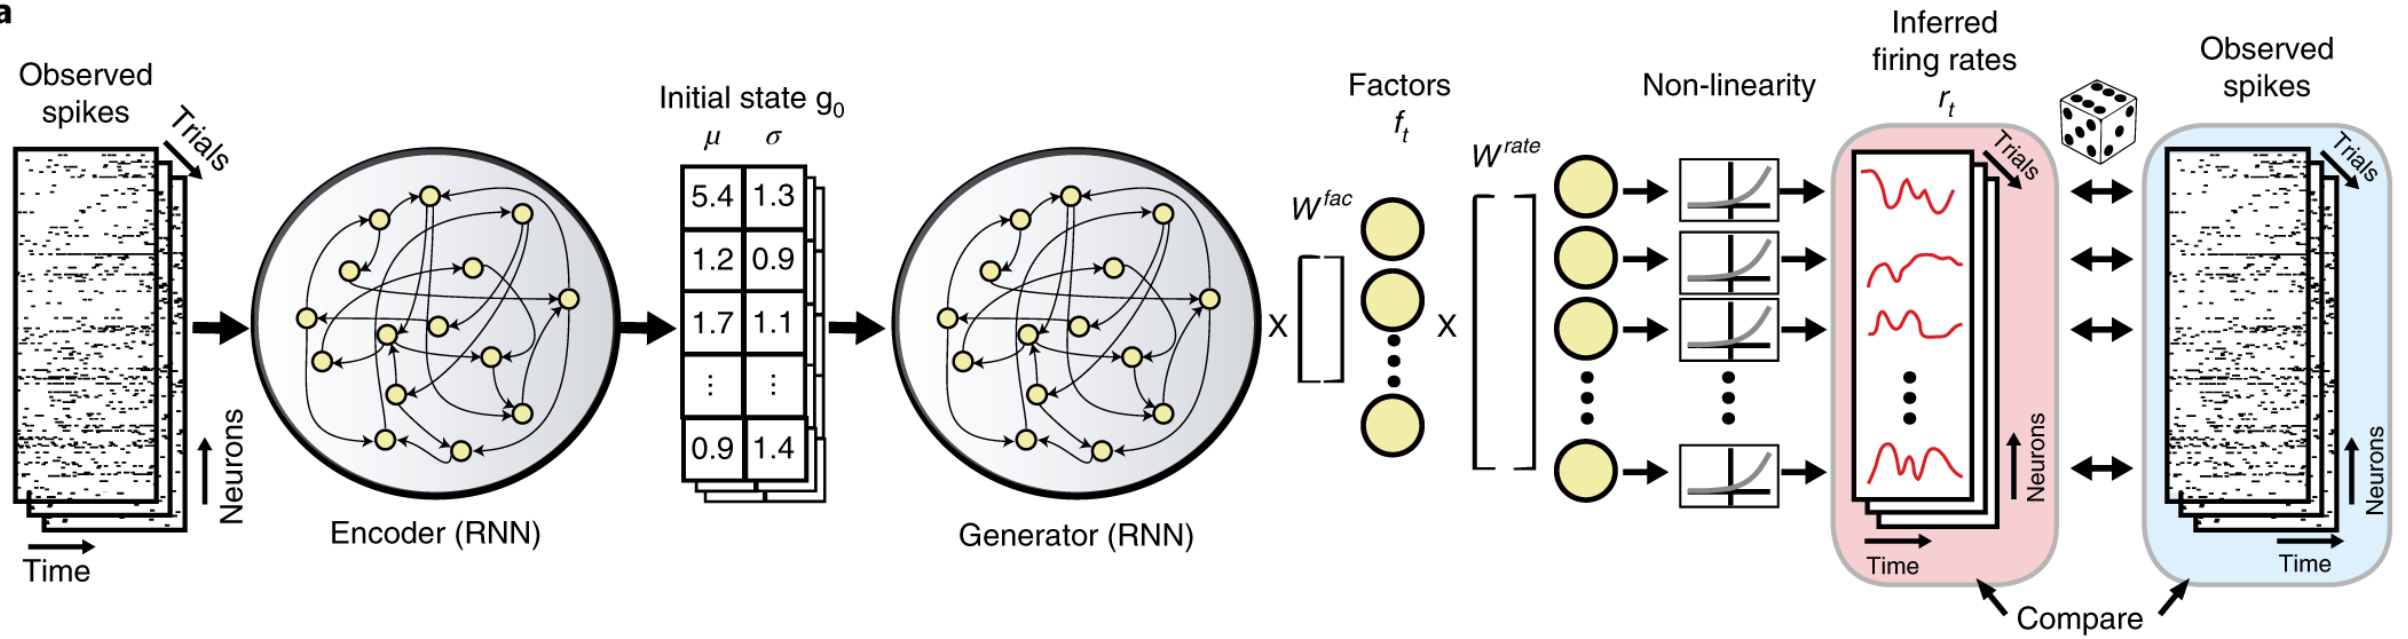
We can compute $p(\mathbf{y}_t | \mathbf{x}_{1:t}) = p(\mathbf{y}_t | \mathbf{z}(\mathbf{x}_{1:t}))$

$$\theta = \underset{\theta}{\operatorname{argmax}} \log p(\mathbf{Y} | \mathbf{X})$$

No more closed-form solutions!

Nonlinear dimensionality reduction


a




Pandarinath et al. (2018); Schimel et al. (2022); Linderman et al. (2016)

The brain as a 'dynamical system'

c Neural flow field



d Initial conditions influence neural trajectory



Additional reading

- Roweis and Ghahramani (1999): *A unifying review of linear Gaussian models*.
 - Excellent of overview of how different linear Gaussian models are related.
- Bishop (2006): *Pattern recognition and machine learning*.
 - Useful introduction to different algorithms used for learning in these and other models.
- Duncker & Sahani (2021): *Dynamics on the manifold: identifying computational dynamical activity from neural population recordings*.
 - A review of different modern approaches used to extract latent dynamical systems and other timeseries models from empirical neural data.
- Vyas et al. (2020): *Computation through neural population dynamics*
 - Review of methods, models and data that model different biological neural circuits as 'dynamical systems'.
- Kao & Hennequin (2019): *Neuroscience out of control: control-theoretic perspectives on neural circuit dynamics*.
 - Introduction to the use of more advanced tools from control theory for understanding neural dynamics.