Worksheet: artificial neural networks

Joaquin Rapela

March 10, 2025

This worksheet is a continuation of the Neuromatch tutorial on decoding neural responses that we covered in the practical session on artificial neural networks.
Here we will:

1. evaluate the performance of the neural network model on test data,

2. use a circular variable loss function,

3. perform model selection (i.e., select optimal number of layers and layer width),

4. visualise the tuning curve of the most relevant neurons,

5. train the model in a GPU.

1 Evaluate the model on test data

Partition the data into a train and a test set. For example, the train set can contain 60% of the total data and the test set the remaining part. Use the train set to estimate the model parameters. For each epoch in the training loop compute the mean-squared error (MSE) on the train set and on the test set. Plot these two errors as a function of iteration, as in Figure 1.

Figure 1: Train and test MSE as a function of iteration. Click on the figure to see its interactive version.

To train the mode you may want to use the script doTrainNet.py with the parameters in the script trainMSETestMSESGD.csh.

To plot the train and test loss function you may want to use the script doPlotTrainTestLosses.py with the parameters in the script plotLossesTrainMSETestMSE.csh.

2 Circular variable loss function

Orientation is a circular variable, but the MSE is not well suited for circular variables. Create a loss function for circular variables. Estimate a model using this new loss function on the train and test data. Build a figure as Fig. 1, but using the circular loss function.

You may want to complete function in function loss, of class CircularLoss, of module myNets.py. Then you can estimate a model with this loss function with the script doTrainNet.py, using parameters train_loss_fn_type=Circular and test_loss_fn_type=Circular. Adjust the parameter learning_rate to obtain best results.

To test which of the MSE or Circular loss is better, estimate a first model with the MSE loss for the train data and the Circular loss for the test data, and a second model with the Circular loss for the train data and the Circular loss for the test data. The model giving the lowest loss in the test data should be the best one. Use the same random seed for both estimations. Is the circular loss better?

3 Model selection

Estimate models with different number of hidden units, compare their test errors, and report the best number of hidden units that you found. Use train and test Circular loss functions
    Optionally:

1. test if using a nonlinear model yields substantially better results than using a linear one.

   Hint: replace `net = myNets.DeepNetReLU(n_neurons, n_hidden)`

   with `net = myNets.DeepNetLinear(n_neurons, n_hidden)`

   in doTrainNet.py.

2. vary the depth of the network.

3. vary the random seed, by setting the parameter `random_seed` of the script doTrainNet.py to different values. If you obtain very different losses with different random seeds, then the loss landscape may have many local minima.

4 Visualise the tuning curve of the most relevant neurons

Not all neurons are equally informative for the decoding task. Neurons with largest weights into the hidden layer should be more relevant. Plot the tuning function of the first 50 neurons with largest weights to the hidden layer, as in Figure ??. Then plot the 50 next neurons with largest weight. See if you observe a pattern in the most relevant cells that is different from the other cells.
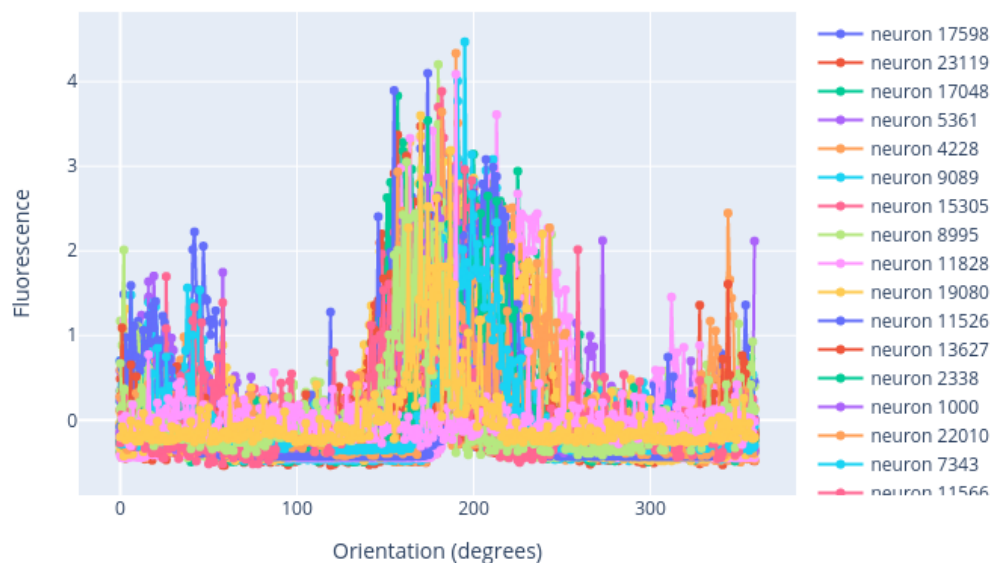

Figure 2: Tuning functions of the 50 neurons with largest weights to the hidden layer. Click on the figure to see its interactive version.

You may want to use the script doPlotTunningCurves.py and set the parameter `rank_1st_neuron`.

5   Train the model in a GPU

If a GPU is available in your system, the script doTrainNet.py will run on the GPU.
At the SWC you can run this script in the cluster by:

1. **ssh into the bastion node**

   `ssh ssh.swc.ucl.ac.uk`

2. **allocate an interactive job in the cluster**

   `srun -p fast -t 2:00:00 --gres=gpu:1 --pty bash -i`

3. **run a script**

   `./trainCircularTestCircularSGD.csh`